

Miniprojeto 2 - Redes Neurais Convolucionais com MNIST

Alunos

- Ítalo Rodrigo Barbosa Paulino - irbp
- José Nilton de Oliveira Lima Júnior - jnolj

O objetivo deste mini projeto é construir um classificador para a base de dados MNIST utilizando uma Rede Neural Convolucional (CNN). Para isso, utilizaremos a biblioteca tensorflow em conjunto com o keras para modelar a nossa rede neural e só assim depois treiná-la e avaliar o seu desempenho com base no conjunto de teste disponibilizado pelo dataset.

Esse relatório seguirá o modelo de código comentado, feito em Python3 e utilizando Jupyter Notebook.

```
In [0]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Flatten, Dense, DepthwiseConv2D
, Conv2D, MaxPool2D, ZeroPadding2D
from tensorflow.keras.layers import BatchNormalization, Activation,
GlobalAveragePooling2D, Dropout
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras.utils import to_categorical

import numpy as np
import matplotlib.pyplot as plt
import os
from sklearn.model_selection import train_test_split
from collections import Counter
from sklearn.metrics import confusion_matrix

from google.colab import files
```

O primeiro passo feito foi carregar o dataset do MNIST e normalizar os valores correspondentes aos seus pixels. Logo após isso transformamos a imagem em um array tridimensional, mas com apenas um channel (a imagem é em tons de cinza), pois será essa o tipo de entrada esperada pela nossa rede neural.

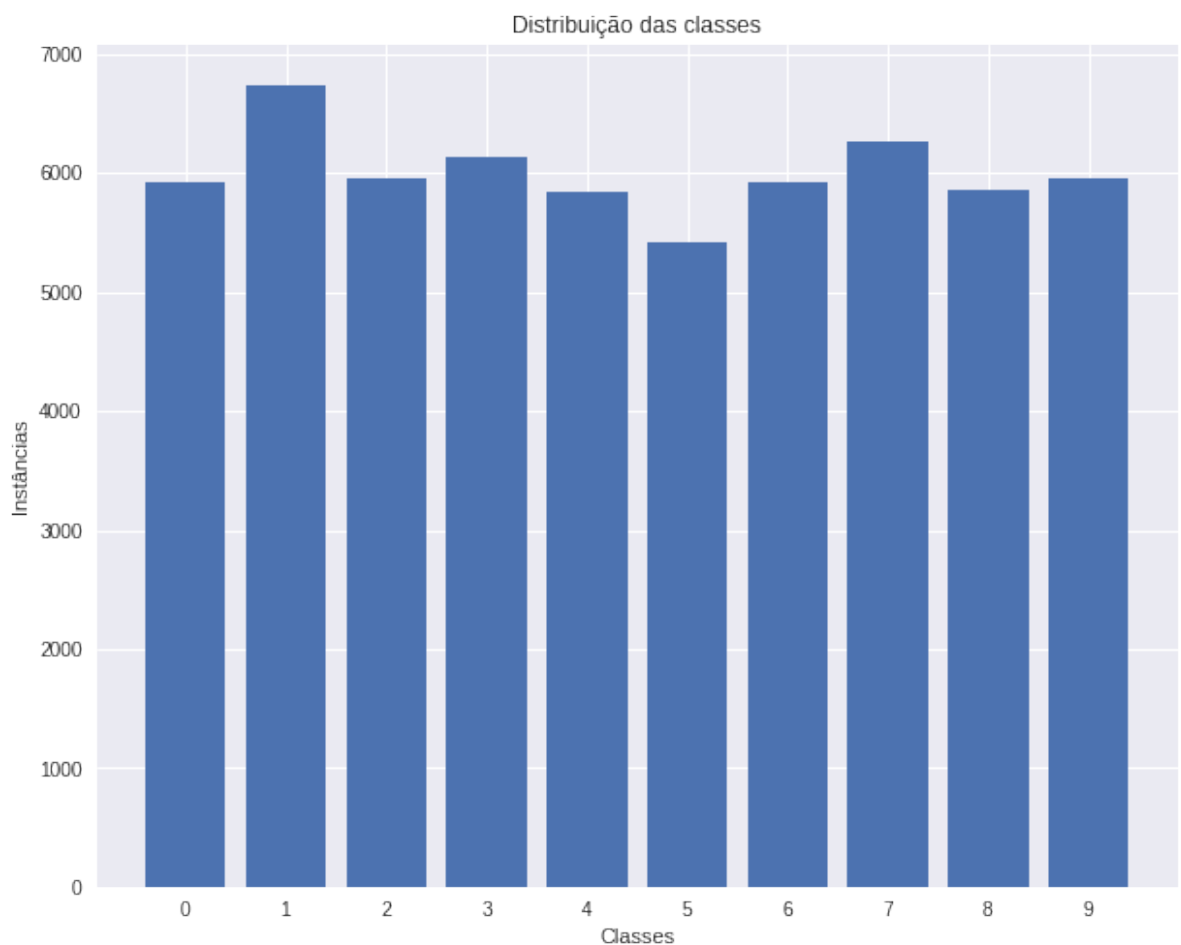
```
In [2]: mnist = keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
x_train, x_test = x_train.reshape(-1, 28, 28, 1), x_test.reshape(-1, 28, 28, 1)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

Out[2]: (60000, 28, 28, 1)

Em seguida plotamos a distribuição de cada classe, ou seja, quantas instâncias de cada classe está presente no nosso dataset.


```
In [3]: unique, count = np.unique(y_train, return_counts=True)
plt.figure(figsize=(10,8))
plt.title("Distribuição das classes")
plt.bar(unique, count)
plt.xticks(unique)
plt.xlabel("Classes")
plt.ylabel("Instâncias")
plt.show()
```



Agora fazemos a divisão do conjunto de treino em conjunto de treino e conjunto de validação que servirá para avaliar o desempenho da nossa rede enquanto a mesma é treinada. Logo depois foi plotado alguns exemplos do nosso dataset.

```
In [0]: y_train, y_test = to_categorical(y_train), to_categorical(y_test)
        random_seed = 2
        x_train, x_val, y_train, y_val = train_test_split(x_train, y_train,
        test_size=0.1, random_state=random_seed)
```

```
In [65]: fig, ax = plt.subplots(5, 5, sharex=True, sharey=True, figsize=(10,
        10))
        index = 0
        for row in range(5):
            for col in range(5):
                ax[row, col].imshow(x_train[index].reshape(28,28))
                ax[row, col].set_title("Label: {}".format(y_train[index].ar
        gmax()))
                ax[row, col].grid(False)
                index += 1
```



Abaixo segue a implementação do nosso modelo de CNN e logo após um sumário mostra a estrutura da rede com o número de parâmetros em cada camada assim como o número total de parâmetros treináveis. Duas camadas de Dropout foram utilizadas para que uma certa proporção de "nós" da nossa rede seja aleatoriamente ignorada, isso força que a rede extraia as features de maneira distribuída ajudando também a reduzir o overfitting.

```
In [0]: model = Sequential()

model.add(Conv2D(filters=32,
                  kernel_size=(5,5),
                  padding="same",
                  activation="relu",
                  input_shape=(28,28,1)))
model.add(Conv2D(filters=32,
                  kernel_size=(5,5),
                  padding="same",
                  activation="relu"))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Dropout(0.25))

model.add(Conv2D(filters=64,
                  kernel_size=(3,3),
                  padding="same",
                  activation="relu"))
model.add(Conv2D(filters=64,
                  kernel_size=(3,3),
                  padding="same",
                  activation="relu"))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(256, activation="relu"))
model.add(Dropout(0.5))
model.add(Dense(10, activation="softmax"))
```

```
In [7]: model.summary()
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 32)	832
conv2d_1 (Conv2D)	(None, 28, 28, 32)	25632
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
dropout (Dropout)	(None, 14, 14, 32)	0
conv2d_2 (Conv2D)	(None, 14, 14, 64)	18496
conv2d_3 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
dropout_1 (Dropout)	(None, 7, 7, 64)	0
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 256)	803072
dropout_2 (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 10)	2570
Total params: 887,530		
Trainable params: 887,530		
Non-trainable params: 0		

```
In [0]: def check_tpu():
        try:
            tpu_address = 'grpc://' + os.environ['COLAB_TPU_ADDR']
            print('TPU address is', tpu_address)

            with tf.Session(tpu_address) as session:
                devices = session.list_devices()

            return True, tpu_address
        except Exception:
            return False, ""
```

```
In [9]: has_tpu, tpu_address = check_tpu()
        if has_tpu:
            tpu_model = tf.contrib.tpu.keras_to_tpu_model(
                model,
                strategy=tf.contrib.tpu.TPUDistributionStrategy(
                    tf.contrib.cluster_resolver.TPUClusterResolver(tpu_address)
                )
            )
```

```

TPU address is grpc://10.79.236.226:8470
INFO:tensorflow:Querying Tensorflow master (b'grpc://10.79.236.226:8470') for TPU system metadata.
INFO:tensorflow:Found TPU system:
INFO:tensorflow:*** Num TPU Cores: 8
INFO:tensorflow:*** Num TPU Workers: 1
INFO:tensorflow:*** Num TPU Cores Per Worker: 8
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:CPU:0, CPU, -1, 10286764627793717633)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_CPU:0, XLA_CPU, 17179869184, 15454925982782434438)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:XLA_GPU:0, XLA_GPU, 17179869184, 3839510994108593343)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:0, TPU, 17179869184, 4868800248739648966)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:1, TPU, 17179869184, 13090091739542284063)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:2, TPU, 17179869184, 9059051702479849442)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:3, TPU, 17179869184, 2881072623438645419)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:4, TPU, 17179869184, 16289019639762527406)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:5, TPU, 17179869184, 13032478147798407067)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:6, TPU, 17179869184, 4236763416648248849)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU:7, TPU, 17179869184, 1496121818056094762)
INFO:tensorflow:*** Available Device: _DeviceAttributes(/job:worker/replica:0/task:0/device:TPU_SYSTEM:0, TPU_SYSTEM, 17179869184, 10225653847187607492)
WARNING:tensorflow:tpu_model (from tensorflow.contrib.tpu.python.tpu.keras_support) is experimental and may change or be removed at any time, and without warning.
INFO:tensorflow:Connecting to: b'grpc://10.79.236.226:8470'

```

Como função otimizadora nós utilizamos o RMSprop, pois ele efetua os ajustes de forma bem simples na tentativa de reduzir a agressividade e diminuindo o learning rate monotonicamente. Poderíamos ter utilizado o Stochastic Gradient Descent (SGD), mas optamos pelo RMSprop por ele apresentar um desempenho melhor.

```
In [0]: optimizer = RMSprop(lr=0.001, rho=0.9, epsilon=1e-08, decay=0.0)
        if has_tpu:
            tpu_model.compile(optimizer=optimizer, loss="categorical_crossentropy", metrics=["accuracy"])
        else:
            model.compile(optimizer=optimizer, loss="categorical_crossentropy", metrics=["accuracy"])
```

```
In [0]: learning_rate_reduction = ReduceLROnPlateau(monitor='val_acc',
                                                    patience=3,
                                                    verbose=1,
                                                    factor=0.5,
                                                    min_lr=0.00001)

epochs = 30
batch_size = 1000
```

```
In [0]: def train_gen(batch_size):
        while True:
            offset = np.random.randint(0, x_train.shape[0] - batch_size)
            yield x_train[offset:offset+batch_size], y_train[offset:offset
+ batch_size]
```

Abaixo segue o histórico de treino para cada uma das 30 épocas que nossa rede foi submetida durante o treino.

```
In [14]: if has_tpu:
        history = tpu_model.fit_generator(train_gen(batch_size),
                                          epochs=epochs,
                                          validation_data=(x_val, y_val
),
                                          steps_per_epoch=x_train.shape
[0]//batch_size,
                                          callbacks=[learning_rate_redu
ction])
    else:
        history = model.fit_generator(train_gen(batch_size),
                                      epochs=epochs,
                                      validation_data=(x_val, y_val
),
                                      steps_per_epoch=x_train.shape
[0]//batch_size,
                                      callbacks=[learning_rate_redu
ction])
```

Epoch 1/30

INFO:tensorflow:New input shapes; (re-)compiling: mode=train, [TensorSpec(shape=(125, 28, 28, 1), dtype=tf.float32, name='conv2d_input0'), TensorSpec(shape=(125, 10), dtype=tf.float32, name='dense_1_target0')]

INFO:tensorflow:Overriding default placeholder.

INFO:tensorflow:Remapping placeholder for conv2d_input

INFO:tensorflow:Cloning RMSprop {'lr': 0.0010000000474974513, 'rho': 0.8999999761581421, 'decay': 0.0, 'epsilon': 1e-08}

```
INFO:tensorflow:Get updates: Tensor("loss/mul:0", shape=(), dtype=
float32)
INFO:tensorflow:Started compiling
INFO:tensorflow:Finished compiling. Time elapsed: 2.26594495773315
43 secs
INFO:tensorflow:Setting weights on TPU model.
53/54 [=====>.] - ETA: 0s - loss: 0.6942 -
acc: 0.7777INFO:tensorflow:New input shapes; (re-)compiling: mode=
eval, [TensorSpec(shape=(125, 28, 28, 1), dtype=tf.float32, name='
conv2d_input0'), TensorSpec(shape=(125, 10), dtype=tf.float32, nam
e='dense_1_target0'))]
INFO:tensorflow:Overriding default placeholder.
INFO:tensorflow:Remapping placeholder for conv2d_input
INFO:tensorflow:Cloning RMSprop {'lr': 0.0010000000474974513, 'rho
': 0.8999999761581421, 'decay': 0.0, 'epsilon': 1e-08}
INFO:tensorflow:Started compiling
INFO:tensorflow:Finished compiling. Time elapsed: 1.30892777442932
13 secs
54/54 [=====] - 8s 149ms/step - loss: 0.6
862 - acc: 0.7804 - val_loss: 0.1524 - val_acc: 0.9520
Epoch 2/30
54/54 [=====] - 2s 34ms/step - loss: 0.15
15 - acc: 0.9538 - val_loss: 0.0340 - val_acc: 0.9907
Epoch 3/30
54/54 [=====] - 2s 34ms/step - loss: 0.09
45 - acc: 0.9702 - val_loss: 0.0215 - val_acc: 0.9947
Epoch 4/30
54/54 [=====] - 2s 35ms/step - loss: 0.05
29 - acc: 0.9831 - val_loss: 0.0111 - val_acc: 0.9947
Epoch 5/30
54/54 [=====] - 2s 37ms/step - loss: 0.05
70 - acc: 0.9830 - val_loss: 0.0168 - val_acc: 0.9960
Epoch 6/30
54/54 [=====] - 2s 37ms/step - loss: 0.04
41 - acc: 0.9877 - val_loss: 0.0149 - val_acc: 0.9960
Epoch 7/30
54/54 [=====] - 2s 39ms/step - loss: 0.03
53 - acc: 0.9886 - val_loss: 0.0092 - val_acc: 0.9973
Epoch 8/30
54/54 [=====] - 2s 37ms/step - loss: 0.03
07 - acc: 0.9901 - val_loss: 0.0116 - val_acc: 0.9973
Epoch 9/30
54/54 [=====] - 2s 38ms/step - loss: 0.03
12 - acc: 0.9904 - val_loss: 0.0118 - val_acc: 0.9973
Epoch 10/30
54/54 [=====] - 2s 36ms/step - loss: 0.02
53 - acc: 0.9926 - val_loss: 0.0044 - val_acc: 0.9987
Epoch 11/30
54/54 [=====] - 2s 34ms/step - loss: 0.01
90 - acc: 0.9941 - val_loss: 0.0064 - val_acc: 0.9987
Epoch 12/30
54/54 [=====] - 2s 38ms/step - loss: 0.02
75 - acc: 0.9917 - val_loss: 0.0100 - val_acc: 0.9973
Epoch 13/30
52/54 [=====>..] - ETA: 0s - loss: 0.0200 -
acc: 0.9945
Epoch 00013: ReduceLROnPlateau reducing learning rate to 0.0005000
```


000237487257.
54/54 [=====] - 2s 38ms/step - loss: 0.01
93 - acc: 0.9947 - val_loss: 0.0048 - val_acc: 0.9987
Epoch 14/30
54/54 [=====] - 2s 42ms/step - loss: 0.02
12 - acc: 0.9930 - val_loss: 0.0057 - val_acc: 0.9973
Epoch 15/30
54/54 [=====] - 2s 42ms/step - loss: 0.01
95 - acc: 0.9936 - val_loss: 0.0082 - val_acc: 0.9973
Epoch 16/30
53/54 [=====>.] - ETA: 0s - loss: 0.0149 -
acc: 0.9962
Epoch 00016: ReduceLROnPlateau reducing learning rate to 0.0002500
000118743628.
54/54 [=====] - 2s 43ms/step - loss: 0.01
47 - acc: 0.9963 - val_loss: 0.0098 - val_acc: 0.9987
Epoch 17/30
54/54 [=====] - 2s 41ms/step - loss: 0.01
38 - acc: 0.9960 - val_loss: 0.0121 - val_acc: 0.9960
Epoch 18/30
54/54 [=====] - 2s 41ms/step - loss: 0.01
66 - acc: 0.9951 - val_loss: 0.0124 - val_acc: 0.9973
Epoch 19/30
52/54 [=====>..] - ETA: 0s - loss: 0.0181 -
acc: 0.9952
Epoch 00019: ReduceLROnPlateau reducing learning rate to 0.0001250
000059371814.
54/54 [=====] - 2s 39ms/step - loss: 0.01
77 - acc: 0.9954 - val_loss: 0.0097 - val_acc: 0.9987
Epoch 20/30
54/54 [=====] - 2s 42ms/step - loss: 0.00
56 - acc: 0.9979 - val_loss: 0.0021 - val_acc: 0.9987
Epoch 21/30
54/54 [=====] - 2s 39ms/step - loss: 0.01
26 - acc: 0.9959 - val_loss: 0.0085 - val_acc: 0.9973
Epoch 22/30
52/54 [=====>..] - ETA: 0s - loss: 0.0107 -
acc: 0.9971
Epoch 00022: ReduceLROnPlateau reducing learning rate to 6.2500002
9685907e-05.
54/54 [=====] - 2s 37ms/step - loss: 0.01
10 - acc: 0.9969 - val_loss: 0.0100 - val_acc: 0.9973
Epoch 23/30
54/54 [=====] - 2s 38ms/step - loss: 0.01
58 - acc: 0.9960 - val_loss: 0.0049 - val_acc: 0.9973
Epoch 24/30
54/54 [=====] - 2s 38ms/step - loss: 0.01
45 - acc: 0.9967 - val_loss: 0.0070 - val_acc: 0.9973
Epoch 25/30
53/54 [=====>.] - ETA: 0s - loss: 0.0127 -
acc: 0.9962
Epoch 00025: ReduceLROnPlateau reducing learning rate to 3.1250001
48429535e-05.
54/54 [=====] - 2s 37ms/step - loss: 0.01
25 - acc: 0.9963 - val_loss: 0.0054 - val_acc: 0.9987
Epoch 26/30
54/54 [=====] - 2s 36ms/step - loss: 0.00

```

86 - acc: 0.9975 - val_loss: 0.0108 - val_acc: 0.9973
Epoch 27/30
54/54 [=====] - 2s 35ms/step - loss: 0.01
06 - acc: 0.9964 - val_loss: 0.0111 - val_acc: 0.9947
Epoch 28/30
53/54 [=====>.] - ETA: 0s - loss: 0.0075 -
acc: 0.9973
Epoch 00028: ReduceLROnPlateau reducing learning rate to 1.5625000
742147677e-05.
54/54 [=====] - 2s 36ms/step - loss: 0.00
85 - acc: 0.9970 - val_loss: 0.0051 - val_acc: 0.9987
Epoch 29/30
54/54 [=====] - 2s 38ms/step - loss: 0.00
60 - acc: 0.9978 - val_loss: 0.0018 - val_acc: 1.0000
Epoch 30/30
54/54 [=====] - 2s 38ms/step - loss: 0.00
80 - acc: 0.9972 - val_loss: 0.0039 - val_acc: 0.9987

```

```

In [70]: if has_tpu:
          loss, acc = tpu_model.evaluate(x_test, y_test)
          cpu_model = tpu_model.sync_to_cpu()
          predictions = cpu_model.predict(x_test)
        else:
          loss, acc = model.evaluate(x_test, y_test)
          predictions = model.predict(x_test)

pred_labels = np.argmax(predictions, axis=1)

```

```


10000/10000 [=====] - 2s 217us/step
INFO:tensorflow:Copying TPU weights to the CPU
Acurácia conjunto de teste: 99.6%

```

Como podemos ver nos gráficos plotados abaixo, a nossa rede obteve um excelente resultado com uma acurácia alta e com uma taxa de erro bastante baixa. Podemos observar também que na maioria das épocas a acurácia do conjunto de validação é maior que a acurácia do conjunto de treino, isso é bastante importante pois indica que não está acontecendo um overfit sobre o conjunto de treino.

```
In [48]: fig, ax = plt.subplots(2, 1, figsize=(10,8))
ax[0].plot(history.history['loss'], color='b', label="Training loss")
ax[0].plot(history.history['val_loss'], color='r', label="Validation loss", axes=ax[0])
legend = ax[0].legend(loc='best', shadow=True)

ax[1].plot(history.history['acc'], color='b', label="Training accuracy")
ax[1].plot(history.history['val_acc'], color='r', label="Validation accuracy", axes=ax[1])
legend = ax[1].legend(loc='best', shadow=True)
```



```

In [0]: def plot_confusion_matrix(cm, classes):
        cmap = plt.get_cmap('Reds')
        tick_marks = classes

        #      cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

        plt.figure(figsize=(10, 10))
        plt.grid(False)
        plt.imshow(cm, interpolation='nearest', cmap=cmap)
        plt.title("Confusion matrix")
        plt.colorbar()
        plt.xticks(tick_marks, classes, rotation=45)
        plt.yticks(tick_marks, classes)

        thresh = cm.max() / 2
        for i in range(cm.shape[0]):
            for j in range(cm.shape[1]):
                plt.text(j, i, format(cm[i, j]),
                        horizontalalignment='center',
                        color='white' if cm[i, j] > thresh else 'black')


        plt.tight_layout()
        plt.ylabel('True label')
        plt.xlabel('Predicted label')

```

```

In [46]: cm = confusion_matrix(y_test.argmax(axis=1), pred_labels)
        classes = np.arange(predictions.shape[1])
        plot_confusion_matrix(cm, classes)

```



```
In [67]: fig, ax = plt.subplots(5, 5, sharex=True, sharey=True, figsize=(10,
10))
index = 0
for row in range(5):
    for col in range(5):
        ax[row, col].imshow(x_test[index].reshape(28,28))
        ax[row, col].set_title("Label: {} \n Predicted: {}".format(y_
test[index].argmax(),
pr
ed_labels[index]))
        ax[row, col].grid(False)
        index += 1
```

